

Practical ML - Week 3

Aura Frizzati

Week 3: Predicting with trees, Random Forests, & Model Based Predictions

Week 3.1: Predicting with trees

Key ideas

- Iteratively split the outcome by splitting the predictive variables into groups
- Evaluate “homogeneity” of the outcome within each group
- Split again if necessary, until you get outcomes that are separated into groups that are **homogeneous enough** or that are **small enough**.

Pros

- Easy to interpret
- Better performance in non-linear settings (in comparison to linear regression models)

Cons

- Without pruning/cross-validation, can lead to overfitting
- Harder to estimate uncertainty (in comparison to linear regression models)
- Results may be variable and depending on the exact values of parameters or the variables that have been collected

Basic algorithm

1. Start with all variables in one group
2. Find the variable/split that best separates the outcome into two different homogenous groups
3. Divide the data into 2 groups ('leaves') on that split ('node')
4. Within each split, find the best variable/split (including variables we have already used to split the groups) that separates the outcome
5. Continue until the groups are too small or sufficiently 'pure' (i.e. homogenous)

Measures of impurity

They are all based on this probability which can be estimated:

$$\hat{p}_{mk} = \frac{1}{N_m} \sum_{x_i \text{ in Leaf } m} (y_i = k)$$

Within a particular group (the m leaf) you have N_m total objects you might consider. You can count the number of times that a particular class $(y_i = k)$ appears in that leaf.

- **Misclassification error:**

$$1 - \hat{p}_{mk(m)}$$

$k(m)$ being the most common k outcome class in the dataset

misclassification error = 0 → perfect purity (no misclassification error)

misclassification error = 0.5 → no purity

- **Gini index:**

$$\sum_{k \neq k'} \hat{p}_{mk} * \hat{p}_{mk'} = \sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk}) = 1 - \sum_{k=1}^K \hat{p}_{mk}^2$$

(to not confuse it with the Gini coefficient used in economics!)

The Gini index is 1 - the sum of the squared probability that a sample belongs to any of the different outcome classes. **Gini index** = 0 → perfect purity (this implies one outcome class has classification probability = 1, while all the others 0)


Gini index = 0.5 → no purity (all of the classes are perfectly balanced within each leaf)

- **Deviance/information gain:**


$$- \sum_{k=1}^K \hat{p}_{mk} \log_2[\hat{p}_{mk}]$$

This measure is called **deviance** if you use \ln or otherwise **information gain** using \log_2 . **Deviance/information gain** = 0 → perfect purity **Deviance/information gain** = 1 → no purity

Measures of impurity



- **Misclassification:** $1/16 = 0.06$
- **Gini:** $1 - [(1/16)^2 + (15/16)^2] = 0.12$
- **Information:**
 $-[1/16 \times \log_2(1/16) + 15/16 \times \log_2(15/16)] = 0.34$



- **Misclassification:** $8/16 = 0.5$
- **Gini:** $1 - [(8/16)^2 + (8/16)^2] = 0.5$
- **Information:**
 $-[1/16 \times \log_2(1/16) + 15/16 \times \log_2(15/16)] = 1$

In the example above, the left panel represents a relatively pure split, while the right one is extremely impure.

Example: Iris data

```
data(iris); library(ggplot2)
names(iris)
```

```
## [1] "Sepal.Length" "Sepal.Width" "Petal.Length" "Petal.Width" "Species"
```

```
table(iris$Species)
```


```
##  
##      setosa versicolor  virginica  
##        50         50         50
```

```
library(caret)  
inTrain<-createDataPartition(y = iris$Species, p=0.7,list=F)  
training<-iris[inTrain,]  
testing<-iris[-inTrain,]  
dim(training);dim(testing)
```

```
## [1] 105  5
```

```
## [1] 45  5
```

```
qplot(Petal.Width, Sepal.Width, colour = Species, data= training)
```



There are 3 very distinct classes, although it could be challenging to predict them for a linear model (but a classification tree would be able to handle non linearity).

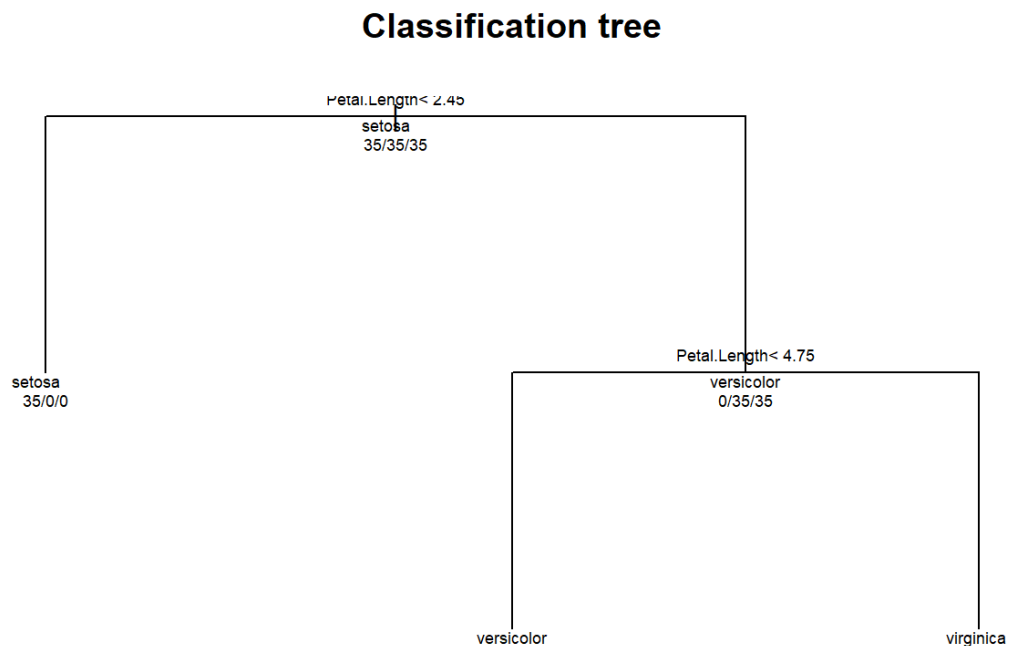
```
modFit<-train(Species ~ .,method = 'rpart', data = training)  
print(modFit$finalModel)
```

```
## n= 105
##
## node), split, n, loss, yval, (yprob)
##      * denotes terminal node
##
## 1) root 105 70 setosa (0.3333333 0.3333333 0.3333333)
##    2) Petal.Length< 2.45 35 0 setosa (1.0000000 0.0000000 0.0000000) *
##    3) Petal.Length>=2.45 70 35 versicolor (0.0000000 0.5000000 0.5000000)
##       6) Petal.Length< 4.75 32 1 versicolor (0.0000000 0.9687500 0.0312500) *
##       7) Petal.Length>=4.75 38 4 virginica (0.0000000 0.1052632 0.8947368) *
```

rpart is an R package for doing regression and classification trees.

`modFit$finalModel` tells you what the final nodes/leaves are, how they split, and the probability for each class to be in each split.

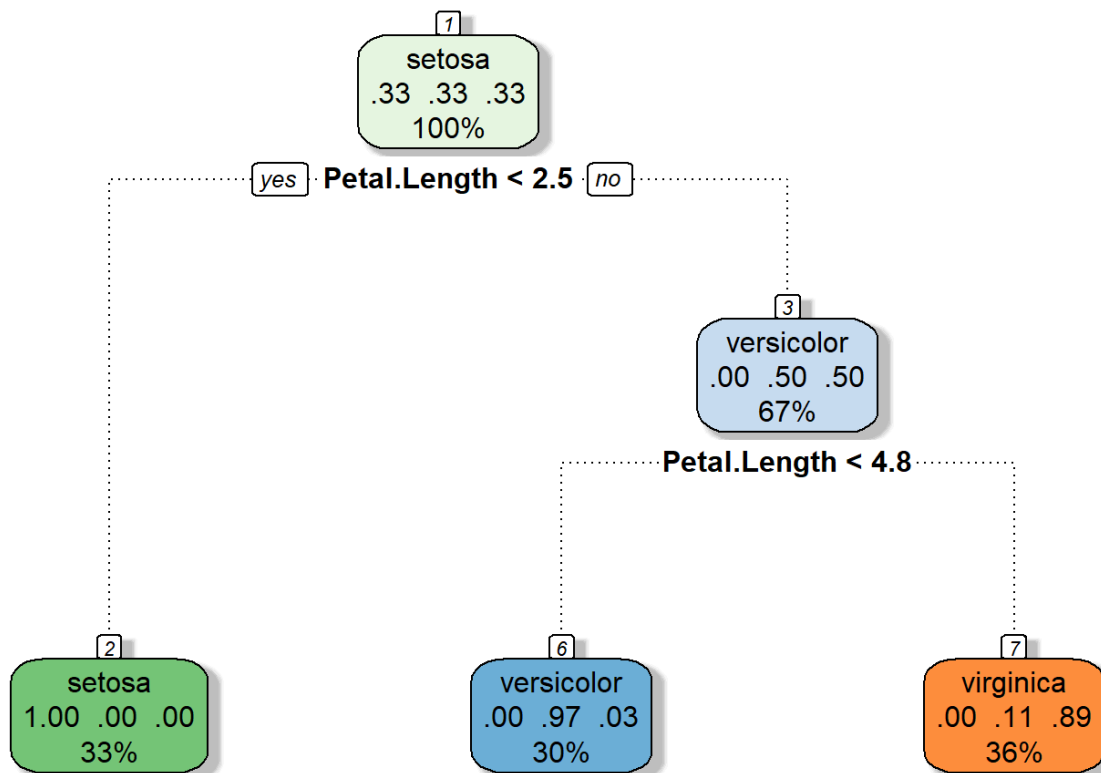
```
plot(modFit$finalModel, uniform = T, main = "Classification tree")
text(modFit$finalModel, use.n = T, all = T, cex =.55)
```



This is a **dendrogram**. The branch to the left denotes when the condition is true, while the branch to the right when it is false.

A prettier version of the same plot can be made with the **rattle** package:

```
library(rattle)
fancyRpartPlot(modFit$finalModel)
```



Rattle 2021-Apr-03 13:49:03 au228742

You can predict new values using the `predict` function:

```
predict(modFit,newdata = testing)
```

```
## [1] setosa setosa setosa setosa setosa setosa
## [7] setosa setosa setosa setosa setosa setosa
## [13] setosa setosa setosa versicolor virginica versicolor
## [19] versicolor versicolor versicolor versicolor virginica versicolor
## [25] versicolor versicolor versicolor versicolor versicolor versicolor
## [31] virginica virginica virginica virginica virginica virginica
## [37] virginica virginica virginica virginica virginica virginica
## [43] virginica virginica virginica
## Levels: setosa versicolor virginica
```

Final notes

- Classification trees are **non-linear models**
 - They use **interactions** between variables
 - Data transformation may be less important (monotone transformations)
 - Trees can also be used for regression problems (continuous outcome)
- Note that there are multiple tree building options in R, both in the caret package (`party` , `rpart`) and out of it (`tree`)

<https://www.amazon.com/Classification-Regression-Trees-Leo-Breiman/dp/0412048418>

(<https://www.amazon.com/Classification-Regression-Trees-Leo-Breiman/dp/0412048418>)

Week 3.2: Bagging

When you fit complicated models, sometimes if you average them together you get a smoother model fit that gives a better balance between potential bias and variance in your fit.

Bootstrap aggregating (bagging)

- Basic idea:
 1. **Resample cases** (similarly to what is done in bootstrapping) and **re-calculate predictions**
 2. **Average** or **majority** of votes from the predictors you have built in this way
- Notes:
 - Similar bias (to the bias you would obtain by fitting any of those models individually)
 - **Reduce variance** (because you have averaged a bunch of predictors together)
 - More useful for **non-linear functions** (e.g. trees or smoothing)

Example on the Ozone data

```
ozone<-read.csv("practical ML course/data/ozone_data.csv", header = T)
ozone<-ozone[order(ozone$ozone),]
head(ozone); dim(ozone)
```

```
##      ozone radiation temperature wind
## 17      1         8          59  9.7
## 19      4        25          61  9.7
## 14      6        78          57 18.4
## 45      7        48          80 14.3
## 106     7        49          69 10.3
## 7       8        19          61 20.1
```

```
## [1] 111  4
```

The objective is to predict 'temperature' as a function of 'ozone'.

Bagged LOESS (Locally Estimated Scatterplot Smoothing)

Local regression or **local polynomial regression**, also known as **moving regression**, is a generalization of moving average and polynomial regression. Its most common methods, initially developed for **scatterplot smoothing**, are **LOESS** (locally estimated scatterplot smoothing) and **LOWESS** (locally weighted scatterplot smoothing). They are two strongly related **non-parametric regression methods** that combine **multiple regression models** in a **k-nearest-neighbor-based meta-model**.

LOESS is a kind of smooth curve that fits through the data. It is similar to a spline model fitting.

```
l1<-matrix(NA,nrow=10,ncol=155)
for(i in 1:10){
  ss<-sample(1:dim(ozone)[1], replace = T) #dim(ozone)[1] = N rows in ozone df
  ozone0<-ozone[ss,] # resampled dataset
  ozone0<-ozone0[order(ozone0$ozone),]
  # the 'span' for loess measures how smooth the fit will be
  loess0<-loess(temperature ~ ozone, data = ozone0, span = 0.2)
  l1[i,]<-predict(loess0,newdata=data.frame(ozone=1:155))
}
head(l1)
```

##	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]
## [1,]	NA	NA	NA	NA	NA	64.89773	70.19492	72.74249
## [2,]	NA	NA	NA	61.21997	62.32309	63.77131	66.03722	70.73670
## [3,]	NA	NA	NA	NA	NA	59.19743	65.10268	69.35389
## [4,]	56.39602	59.63955	62.63709	65.31734	67.87637	70.38038	72.56327	74.15892
## [5,]	NA	NA	NA	NA	NA	60.13440	67.21734	71.16277
## [6,]	59.00478	60.67643	62.22088	63.63669	64.92240	66.07655	67.09770	67.93454
##	[,9]	[,10]	[,11]	[,12]	[,13]	[,14]	[,15]	[,16]
## [1,]	73.25019	73.30859	74.73348	73.35884	70.50015	69.60005	71.99031	72.77043
## [2,]	73.22719	71.54321	68.58573	69.19598	71.17019	71.52438	70.93746	70.01851
## [3,]	71.16744	73.11886	73.83053	73.07516	70.67103	65.90210	71.03505	76.33333
## [4,]	74.90121	73.17130	69.54214	67.03530	66.96636	69.67536	74.03955	77.16881
## [5,]	72.77666	72.07993	70.30595	69.95828	68.41948	68.22835	73.25946	77.00663
## [6,]	68.32931	68.96017	69.55882	70.23541	73.71745	71.99616	72.84947	72.78788
##	[,17]	[,18]	[,19]	[,20]	[,21]	[,22]	[,23]	[,24]
## [1,]	66.58413	60.75302	64.51590	73.20276	78.06522	75.39655	72.53654	71.58738
## [2,]	68.14875	66.84219	69.84663	74.63742	76.30170	75.63740	72.47871	73.67577
## [3,]	71.73777	66.09308	67.52610	70.96390	73.14435	73.14302	72.68402	72.64444
## [4,]	74.98300	72.95066	73.65476	75.88432	77.69295	75.96519	74.20216	74.93017
## [5,]	72.38359	66.29882	66.42095	70.70411	76.50000	76.84167	74.25000	69.66667
## [6,]	67.45751	63.25643	68.88436	73.85041	76.14286	74.48647	72.45455	72.65237
##	[,25]	[,26]	[,27]	[,28]	[,29]	[,30]	[,31]	[,32]
## [1,]	73.03001	75.96026	78.59466	79.14975	77.02969	73.18822	71.93484	71.35385
## [2,]	76.70826	79.82081	81.25800	79.73113	77.92801	76.18355	74.79063	74.53118
## [3,]	74.65186	77.71257	79.62279	80.21347	79.37657	76.68228	73.77261	72.08506
## [4,]	76.31682	77.82801	78.92964	79.08760	77.54382	75.97492	74.00341	72.98444
## [5,]	70.52887	77.90490	82.47890	81.81228	80.53134	79.91751	79.16904	78.05240
## [6,]	73.58947	75.02805	76.73032	78.45848	79.97475	81.04131	81.42038	81.06084
##	[,33]	[,34]	[,35]	[,36]	[,37]	[,38]	[,39]	[,40]
## [1,]	71.63348	72.96197	76.07344	79.13338	81.33347	81.79786	81.99047	81.83223
## [2,]	74.51815	74.73486	75.16461	75.56737	76.79215	78.25512	79.44110	80.39075
## [3,]	70.87547	70.29807	70.48626	71.28081	72.39707	73.55037	74.45606	75.29567
## [4,]	72.82301	73.52076	74.66121	75.82785	77.25688	78.97567	80.25456	80.68209
## [5,]	77.28805	76.38872	75.67667	75.47412	76.00601	77.07004	78.34762	79.52015
## [6,]	80.22689	79.15368	78.07639	77.23018	76.85022	76.93353	77.24666	77.67907
##	[,41]	[,42]	[,43]	[,44]	[,45]	[,46]	[,47]	[,48]
## [1,]	81.37632	80.96667	80.27018	79.61736	79.33872	79.76479	80.68953	81.59856
## [2,]	81.14473	81.74370	82.22834	82.63930	83.01724	83.40284	83.74790	83.98228
## [3,]	76.19797	77.02623	77.88023	78.75087	79.62902	80.63437	81.68147	82.45355
## [4,]	80.61293	79.80382	78.93231	78.67597	79.24299	79.91267	80.82915	81.91759
## [5,]	80.28737	80.70640	80.99475	81.36991	81.68418	81.96346	82.22544	82.54504
## [6,]	78.12019	78.68375	79.44774	80.24805	81.39089	82.81376	83.79949	84.32025
##	[,49]	[,50]	[,51]	[,52]	[,53]	[,54]	[,55]	[,56]
## [1,]	82.35983	82.84128	83.09234	83.26101	83.36316	83.41467	83.43139	83.42920
## [2,]	84.11366	84.14973	84.04338	83.77445	83.39109	82.94143	82.47363	82.03582
## [3,]	82.95696	83.40037	83.77056	84.05429	84.27791	84.47866	84.65972	84.82425
## [4,]	83.10318	84.21556	85.09263	85.67187	85.87461	85.72607	85.31959	84.74849
## [5,]	82.57831	82.63414	82.58329	82.44850	82.25245	82.01788	81.76748	81.52398
## [6,]	84.70330	85.20516	85.15499	84.76440	84.12581	83.33162	82.47423	81.64606
##	[,57]	[,58]	[,59]	[,60]	[,61]	[,62]	[,63]	[,64]
## [1,]	83.42396	83.43154	83.46782	83.54864	83.68989	83.90743	84.21713	84.67236
## [2,]	81.67616	81.44278	81.38383	81.47413	81.63391	81.82941	82.02688	82.19257
## [3,]	84.97542	85.11640	85.25037	85.38048	85.50991	85.64183	85.77941	85.92581
## [4,]	84.10611	83.48580	82.98087	82.68468	82.69055	83.02864	83.56769	84.15893
## [5,]	81.31007	81.14847	81.06190	81.07306	81.20466	81.47942	81.92004	82.58392
## [6,]	80.93951	80.44699	80.26090	80.42279	80.85741	81.48200	82.21379	82.97003
##	[,65]	[,66]	[,67]	[,68]	[,69]	[,70]	[,71]	[,72]
## [1,]	85.27968	85.98476	86.73325	87.47081	88.14310	88.69578	89.07451	89.16987
## [2,]	82.29204	82.49738	82.77561	83.11651	83.50982	83.94532	84.41276	84.90190

```


## [3,] 86.08420 86.25775 86.44963 86.66302 86.90107 87.16696 87.46386 87.80936
## [4,] 84.65355 85.16691 85.85839 86.64527 87.44483 88.17435 88.75112 89.19972
## [5,] 83.47560 84.53851 85.71608 86.95174 88.18894 89.37110 90.44165 91.34404
## [6,] 83.66795 84.45818 85.46159 86.55076 87.59831 88.47684 89.05894 89.27983
##      [,73]  [,74]  [,75]  [,76]  [,77]  [,78]  [,79]  [,80]
## [1,] 89.24150 89.39903 89.47479 89.55201 89.41871 89.08306 88.95853 89.17272
## [2,] 85.40250 85.90433 86.39713 86.87068 87.31473 87.71904 88.07337 88.41963
## [3,] 88.20535 88.63141 89.06714 89.49212 89.88594 90.22818 90.49845 90.73277
## [4,] 89.61742 90.00436 90.34588 90.64600 90.90872 91.13808 91.33808 91.47046
## [5,] 92.02168 92.46930 92.79870 93.12896 93.45577 93.70187 93.87687 93.99039
## [6,] 89.44954 89.64171 89.71112 89.76436 89.90803 90.09349 90.23861 90.39974
##      [,81]  [,82]  [,83]  [,84]  [,85]  [,86]  [,87]  [,88]
## [1,] 89.38645 89.67172 90.00443 90.36042 90.71557 91.04577 91.32686 91.53473
## [2,] 88.79012 89.16021 89.50528 89.80068 90.02179 90.21561 90.43125 90.64723
## [3,] 90.96900 91.18943 91.37636 91.51208 91.61332 91.71173 91.81928 91.89858
## [4,] 91.51764 91.51467 91.49659 91.49847 91.55534 91.67014 91.80271 91.92351
## [5,] 94.05202 94.07138 94.05809 94.02175 93.97198 93.90347 93.81026 93.70101
## [6,] 90.57096 90.71850 90.85693 91.00082 91.16474 91.31337 91.41595 91.49391
##      [,89]  [,90]  [,91]  [,92]  [,93]  [,94]  [,95]  [,96]
## [1,] 91.64524 91.67955 91.67823 91.64631 91.58883 91.51081 91.41727 91.31326
## [2,] 90.84204 90.99421 91.08223 91.09335 91.03463 90.91433 90.74073 90.52211
## [3,] 91.91223 91.83870 91.69472 91.50030 91.27548 91.04027 90.81469 90.61878
## [4,] 92.00300 92.07154 92.16776 92.27452 92.37467 92.45109 92.48662 92.46412
## [5,] 93.58438 93.46901 93.36356 93.25743 93.13479 92.99646 92.84324 92.67594
## [6,] 91.56870 91.64583 91.71688 91.78237 91.84285 91.89884 91.95090 91.99954
##      [,97]  [,98]  [,99]  [,100]  [,101]  [,102]  [,103]  [,104]
## [1,] 91.18085 91.00490 90.79451 90.55881 90.30694 90.04801 89.79115 89.54549
## [2,] 90.26671 89.98283 89.67873 89.36268 89.04294 88.72779 88.42551 88.14435
## [3,] 90.40869 90.13573 89.81262 89.45207 89.06679 88.66950 88.27293 87.88977
## [4,] 92.36646 92.20715 92.01379 91.79040 91.54099 91.26956 90.98013 90.67671
## [5,] 92.49537 92.30235 92.09768 91.88218 91.65666 91.42192 91.17878 90.92805
## [6,] 92.04530 92.08873 92.13035 92.17070 92.21031 92.24973 92.28947 92.33009
##      [,105]  [,106]  [,107]  [,108]  [,109]  [,110]  [,111]  [,112]
## [1,] 89.32015 89.12426 88.96694 88.85732 88.77852 88.70758 88.64449 88.58924
## [2,] 87.89259 87.67851 87.51036 87.39642 87.30587 87.20685 87.10607 87.01025
## [3,] 87.53276 87.21461 86.94802 86.74573 86.57861 86.41054 86.24372 86.08033
## [4,] 90.36331 90.04394 89.72261 89.40334 89.09012 88.78698 88.49792 88.22695
## [5,] 90.67054 90.40706 90.13843 89.86544 89.58892 89.30967 89.02850 88.74623
## [6,] 92.37212 92.41608 92.46252 92.51197 92.56496 92.62204 92.68373 92.75057
##      [,113]  [,114]  [,115]  [,116]  [,117]  [,118]  [,119]  [,120]
## [1,] 88.54182 88.50222 88.47044 88.44647 88.43029 88.42190 88.42128 88.42844
## [2,] 86.92610 86.86034 86.81967 86.80298 86.80364 86.82009 86.85081 86.89423
## [3,] 85.92256 85.77262 85.63269 85.50497 85.39165 85.29492 85.20673 85.11764
## [4,] 87.97808 87.75533 87.56271 87.38386 87.20108 87.01635 86.83166 86.64900
## [5,] 88.46366 88.18161 87.90089 87.62231 87.34667 87.07480 86.80749 86.54557
## [6,] 92.82309 92.90184 92.98734 93.08013 93.18075 93.28973 93.39250 93.47443
##      [,121]  [,122]  [,123]  [,124]  [,125]  [,126]  [,127]  [,128]
## [1,] 88.44336 88.46602      NA      NA      NA      NA      NA      NA
## [2,] 86.94884 87.01307      NA      NA      NA      NA      NA      NA
## [3,] 85.02772 84.93705 84.84572 84.75379 84.66136 84.56849 84.47528 84.38180
## [4,] 86.47034 86.29768 86.12928 85.96212 85.79623 85.63167 85.46846 85.30666
## [5,] 86.28983 86.04110 85.80019 85.56789 85.34503 85.13242 84.93086 84.74117
## [6,] 93.53593 93.57741 93.59929 93.60199 93.58592 93.55149 93.49913 93.42925
##      [,129]  [,130]  [,131]  [,132]  [,133]  [,134]  [,135]  [,136]
## [1,]      NA      NA      NA      NA      NA      NA      NA      NA
## [2,]      NA      NA      NA      NA      NA      NA      NA      NA
## [3,] 84.28813 84.19435 84.10054 84.00679 83.91316 83.81974 83.72661 83.63386
## [4,] 85.14631 84.98745 84.83011 84.67435 84.52020 84.36771 84.21691      NA
## [5,] 84.56416 84.40063 84.25140 84.11729 83.99909 83.89763 83.81370      NA

```



```
## [6,] 93.34226 93.23857 93.11862 92.98280 92.83154 92.66525 92.48435 92.28925
##      [,137] [,138] [,139] [,140] [,141] [,142] [,143] [,144]
## [1,]      NA      NA      NA      NA      NA      NA      NA      NA
## [2,]      NA      NA      NA      NA      NA      NA      NA      NA
## [3,] 83.54155 83.44977 83.35860 83.26812 83.17841 83.08955 83.00162 82.9147
## [4,]      NA      NA      NA      NA      NA      NA      NA      NA
## [5,]      NA      NA      NA      NA      NA      NA      NA      NA
## [6,] 92.08036 91.85812 91.62292 91.37518 91.11533 90.84377 90.56092 90.2672
##      [,145] [,146] [,147] [,148] [,149] [,150] [,151] [,152]
## [1,]      NA      NA      NA      NA      NA      NA      NA      NA
## [2,]      NA      NA      NA      NA      NA      NA      NA      NA
## [3,] 82.82887 82.74421 82.66079 82.57871 82.49803 82.41884 82.34122 82.26524
## [4,]      NA      NA      NA      NA      NA      NA      NA      NA
## [5,]      NA      NA      NA      NA      NA      NA      NA      NA
## [6,] 89.96303 89.64881 89.32496 88.99191 88.65006 88.29983 87.94163 87.57589
##      [,153] [,154] [,155]
## [1,]      NA      NA      NA
## [2,]      NA      NA      NA
## [3,] 82.19100 82.11856 82.04801
## [4,]      NA      NA      NA
## [5,]      NA      NA      NA
## [6,] 87.20302 86.82343 86.43753
```

```
plot(ozone$ozone, ozone$temperature, pch = 19, cex = 0.5)
for(i in 1:10){lines(1:155, ll[i,], col='grey', lwd=2)}
lines(1:155, apply(ll, 2, mean), col='red', lwd=2)
```



- In the plot,

each grey line represents the LOESS fit with one resampled dataset. - The grey lines have a lot of curviness and possibly overfit the dataset variability. - The red line is the average of all the fitted grey curves.

There is a proof that shows the **bagging estimate** will always have **lower variability** but **similar bias** to each of the **individual model fits** from which it has been created.

Bagging in caret


There are some models that automatically perform bagging for you → in the `train` function consider these `method` options: - `bagEarth` - `treebag` - `bagFDA`

Alternatively, you can bag any model you choose using the `bag` function.

Example of custom bagging

```
predictors = data.frame(ozone=ozone$ozone) ## predictor/s
temperature = ozone$temperature ## outcome variable
treebag<-bag(predictors, temperature, B=10, ## B = n of subsamples to take

# wit bagControl we specify how to fit the model:
bagControl = bagControl(fit = ctreeBag$fit, ## function to be applied to train the model using each replicated subsample (the train function of the caret package can be called here as well)
predict = ctreeBag$pred, ## call the predict function from the train model
aggregate = ctreeBag$aggregate) ## it specifies how to put all predictions together (e.g. take average)
)
plot(ozone$ozone, temperature, col = 'lightgrey', pch =19)
points(ozone$ozone, predict(treebag$fits[[1]]$fit, predictors), pch =19, col='red')
points(ozone$ozone, predict(treebag, predictors), pch =19, col='blue')
```



represent actual data points - The red dots represent the fit from a single conditional regression tree (it does not capture the data trend very well) - The blue dots represent the fit from the bag regression (i.e. average of the single regression predictions) → data trends are captured better!

Parts of bagging

```
ctreeBag$fit
```

```
## function (x, y, ...)  
## {  
##   loadNamespace("party")  
##   data <- as.data.frame(x, stringsAsFactors = TRUE)  
##   data$y <- y  
##   party::ctree(y ~ ., data = data)  
## }  
## <bytecode: 0x00000000278cf198>  
## <environment: namespace:caret>
```

- The `ctreeBag$fit` function takes the predictor `df (x)` and the predictor (`y`) and calls the function `ctree` to train a **conditional regression tree** on the dataset. The model fit is returned.

```
ctreeBag$pred
```

```
## function (object, x)  
## {  
##   if (!is.data.frame(x))  
##     x <- as.data.frame(x, stringsAsFactors = TRUE)  
##   obsLevels <- levels(object@data@get("response")[, 1])  
##   if (!is.null(obsLevels)) {  
##     rawProbs <- party::treeresponse(object, x)  
##     probMatrix <- matrix(unlist(rawProbs), ncol = length(obsLevels),  
##       byrow = TRUE)  
##     out <- data.frame(probMatrix)  
##     colnames(out) <- obsLevels  
##     rownames(out) <- NULL  
##   }  
##   else out <- unlist(party::treeresponse(object, x))  
##   out  
## }  
## <bytecode: 0x00000000278cfbe0>  
## <environment: namespace:caret>
```

- The `ctreeBag$pred` function takes as input the object created by the `ctreeBag$fit` . The new outcomes from the object and data input are calculated using `treeresponse` .
- These predicted values are then used by the next aggregation function and pooled together (in the example, the median prediction is used as pooled outcome prediction):

```
ctreeBag$aggregate
```

```
## function (x, type = "class")
## {
##   if (is.matrix(x[[1]]) | is.data.frame(x[[1]])) {
##     pooled <- x[[1]] & NA
##     classes <- colnames(pooled)
##     for (i in 1:ncol(pooled)) {
##       tmp <- lapply(x, function(y, col) y[, col], col = i)
##       tmp <- do.call("rbind", tmp)
##       pooled[, i] <- apply(tmp, 2, median)
##     }
##     if (type == "class") {
##       out <- factor(classes[apply(pooled, 1, which.max)],
##         levels = classes)
##     }
##     else out <- as.data.frame(pooled, stringsAsFactors = TRUE)
##   }
##   else {
##     x <- matrix(unlist(x), ncol = length(x))
##     out <- apply(x, 1, median)
##   }
##   out
## }
## <bytecode: 0x0000000278cb4f0>
## <environment: namespace:caret>
```

- Bagging is most useful for nonlinear models.
- It is often used with trees. One of its extension is **random forests**.
- Several models use bagging in caret's `train` function.

https://stat.ethz.ch/education/semesters/FS_2008/CompStat/sk-ch8.pdf
[\(https://stat.ethz.ch/education/semesters/FS_2008/CompStat/sk-ch8.pdf\)](https://stat.ethz.ch/education/semesters/FS_2008/CompStat/sk-ch8.pdf)

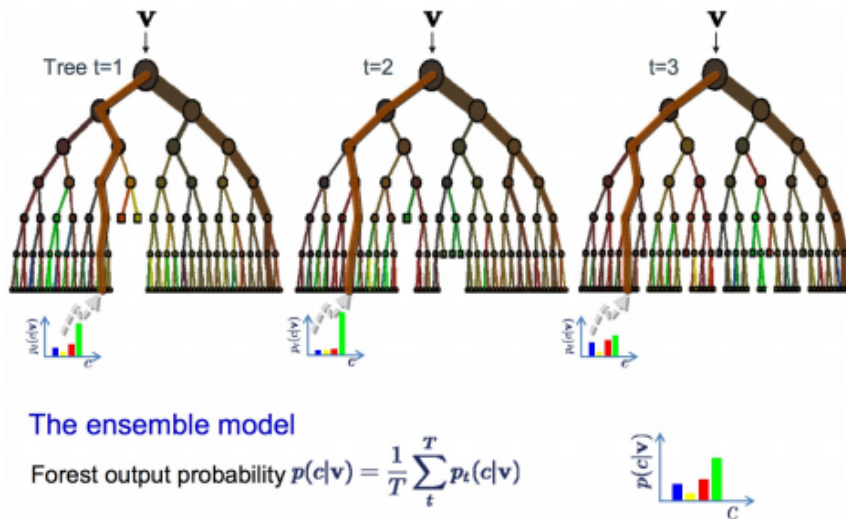
Week 3.3: Random forests

They can be defined as an extension to bagging for classification and regression trees.

1. **Bootstrap samples**
2. **At each split, bootstrap variables** (thus, only a subset of the variables is considered at each potential split. In this way, a diverse set of potential trees can be built)
3. **Grow multiple** (large number) **trees** and vote (or take the average of predictions)

Pros: - **Accuracy**

Cons: - **Low speed** (slow as it needs to build a large number of trees) - ***Poor interpretability** (as the model relies on **bootstrapped samples** with **bootstrapped nodes**) - **Overfitting** (hard to understand which trees are leading to overfitting → very important to use **crossvalidation**)



- The basic idea is that you build a large number of trees (T), each created from a bootstrap sample.
- At each node of one tree we allow a subsample of the variables to potentially contribute to the splits.
- The same observation (v in the figure above) will end up at possibly a different leaf at the bottom of each tree, corresponding to a particular prediction.
- All the different predictions from all the prediction trees are then averaged (**ensemble model**) to create the final class (c) prediction for the observation (v).

Iris data example

```
data(iris);library(ggplot2);library(caret);library(randomForest)
inTrain<-createDataPartition(y=iris$Species, p=0.7,list=F)
training<-iris[inTrain,]
testing<-iris[-inTrain,]
summary(training)
```

```
##   Sepal.Length   Sepal.Width   Petal.Length   Petal.Width
##   Min.    :4.300   Min.    :2.000   Min.    :1.100   Min.    :0.100
##   1st Qu.:5.100   1st Qu.:2.800   1st Qu.:1.500   1st Qu.:0.300
##   Median :5.800   Median :3.000   Median :4.400   Median :1.300
##   Mean    :5.823   Mean    :3.055   Mean    :3.749   Mean    :1.209
##   3rd Qu.:6.400   3rd Qu.:3.300   3rd Qu.:5.100   3rd Qu.:1.800
##   Max.    :7.900   Max.    :4.400   Max.    :6.900   Max.    :2.500
##           Species
##   setosa    :35
##   versicolor:35
##   virginica :35
##
##
##
```

```
### RANDOM FORESTS
modFit<-train(data=training,Species ~ .,method='rf',prox=T)# 'rf' = random forest,
# prox=T is to allow to visualise class centers (see below)
modFit
```

```
## Random Forest
##
## 105 samples
## 4 predictor
## 3 classes: 'setosa', 'versicolor', 'virginica'
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 105, 105, 105, 105, 105, 105, ...
## Resampling results across tuning parameters:
##
## mtry Accuracy Kappa
## 2 0.9533341 0.929197
## 3 0.9533341 0.929197
## 4 0.9533341 0.929197
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 2.
```

A bunch of different tuning parameters were tried during training → `mtry` = Number of variables randomly sampled as candidates at each split

```
getTree(modFit$finalModel,k=2) # k specifies which tree we want to examine (number 2 in this case)
```

```
## left daughter right daughter split var split point status prediction
## 1 2 3 1 5.45 1 0
## 2 4 5 4 0.75 1 0
## 3 6 7 4 1.75 1 0
## 4 0 0 0 0.00 -1 1
## 5 8 9 3 4.20 1 0
## 6 10 11 4 0.60 1 0
## 7 12 13 2 3.15 1 0
## 8 0 0 0 0.00 -1 2
## 9 0 0 0 0.00 -1 3
## 10 0 0 0 0.00 -1 1
## 11 0 0 0 0.00 -1 2
## 12 0 0 0 0.00 -1 3
## 13 14 15 3 5.05 1 0
## 14 0 0 0 0.00 -1 2
## 15 0 0 0 0.00 -1 3
```

- Each row corresponds to a split
- “var split” is the variable that was used for the split
- “split point” is the value of the variable used for the split


Class ‘centers’

- They specify the centers of the class predictions:

```
library(randomForest);library(ggplot2)
#training[,c(3,4)] denote the two predictors to represent in the plot:
irisP<-classCenter(training[,c(3,4)],training$Species, modFit$finalModel$prox) ##prox is used here
irisP<-as.data.frame(irisP);irisP$Species<-rownames(irisP)
irisP
```

```
##          Petal.Length Petal.Width  Species
## setosa          1.4         0.2    setosa
## versicolor      4.4         1.3 versicolor
## virginica       5.5         2.0  virginica
```

```
p<-qplot(Petal.Width,Petal.Length,col=Species,data=training)
p + geom_point(aes(x=Petal.Width,y=Petal.Length,col=Species), size=5,shape=4,data=irisP)
```




Predicting new values

```
library(caret)
pred<-predict(modFit,testing)
testing$predRight<-pred == testing$Species
table(pred,testing$Species)
```

```
##
## pred      setosa versicolor virginica
## setosa    15      0          0
## versicolor 0      15          3
## virginica  0      0          12
```

```
qplot(Petal.Width,Petal.Length,colour=predRight, data=testing, main = 'New data prediction')
```

New data prediction



```
confusionMatrix(pred,testing$Species)
```



```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  setosa versicolor virginica
##   setosa      15          0          0
##   versicolor   0         15          3
##   virginica    0          0         12
##
## Overall Statistics
##
##           Accuracy : 0.9333
##           95% CI : (0.8173, 0.986)
##   No Information Rate : 0.3333
##   P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9
##
##   Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: setosa Class: versicolor Class: virginica
## Sensitivity           1.0000           1.0000           0.8000
## Specificity           1.0000           0.9000           1.0000
## Pos Pred Value        1.0000           0.8333           1.0000
## Neg Pred Value        1.0000           1.0000           0.9091
## Prevalence            0.3333           0.3333           0.3333
## Detection Rate        0.3333           0.3333           0.2667
## Detection Prevalence  0.3333           0.4000           0.2667
## Balanced Accuracy      1.0000           0.9500           0.9000
```

- The two misclassified samples were at the border between 2 classes.

Notes and further resources

- Random forests are usually one of the top performing algorithms along with boosting in prediction contests.
- Random forests are **difficult to interpret** but **very accurate**.
- Care should be taken to avoid overfitting, using cross-validation with the `rf` library (see `rfcv` function: <https://cran.r-project.org/web/packages/randomForest/randomForest.pdf> (<https://cran.r-project.org/web/packages/randomForest/randomForest.pdf>)) or indirectly via `caret`.
- Website of creator of random forests: https://www.stat.berkeley.edu/~breiman/RandomForests/cc_home.htm (https://www.stat.berkeley.edu/~breiman/RandomForests/cc_home.htm)

Week 3.4: Boosting

Key ideas

1. Take lots of (possibly) weak predictors
2. Weight them (to take advantage of their strength) and add them up
3. Get a stronger predictor

Basic idea behind boosting

1. Start with a **set of classifiers** h_1, \dots, h_k . They are usually from the same class of classifiers (examples: all possible trees, all possible regression models, all possible cutoffs).

2. Create a classifier that **combines classification functions**: $f(x) = \text{sgn} \left(\sum_{t=1}^T a_t h_t(x) \right)$ (a_t is a weight, $h_t(x)$ is a classifier)


- Goal is to **minimise error** (on training set)
- **Iterative**: select one h at each step
- Calculate **weights based on errors**
- Upweight missed classifications and select next h

The most famous boosting algorithm is **Ada boosting**.

Simple example

<https://alliance.seas.upenn.edu/~cis520/dynamic/2020/wiki/index.php?n=Lectures.Boosting>


(<https://alliance.seas.upenn.edu/~cis520/dynamic/2020/wiki/index.php?n=Lectures.Boosting>)



- We are trying to separate the +/blue category from the -/red category using 2 dimensions/predictors (X- and Y-axes)

Round 1: adaboost


Round 1




- The first classifier considers anything to the left of the red line as a +/blue and anything to the right as a -/red.
- 3 points were misclassified
- ϵ_1 is the error rate
- The weight of the misclassified samples is increased (**upweighting**), so they get more importance for building the next classifier.

Round 2 & 3: adaboost

Round 2



Round 3

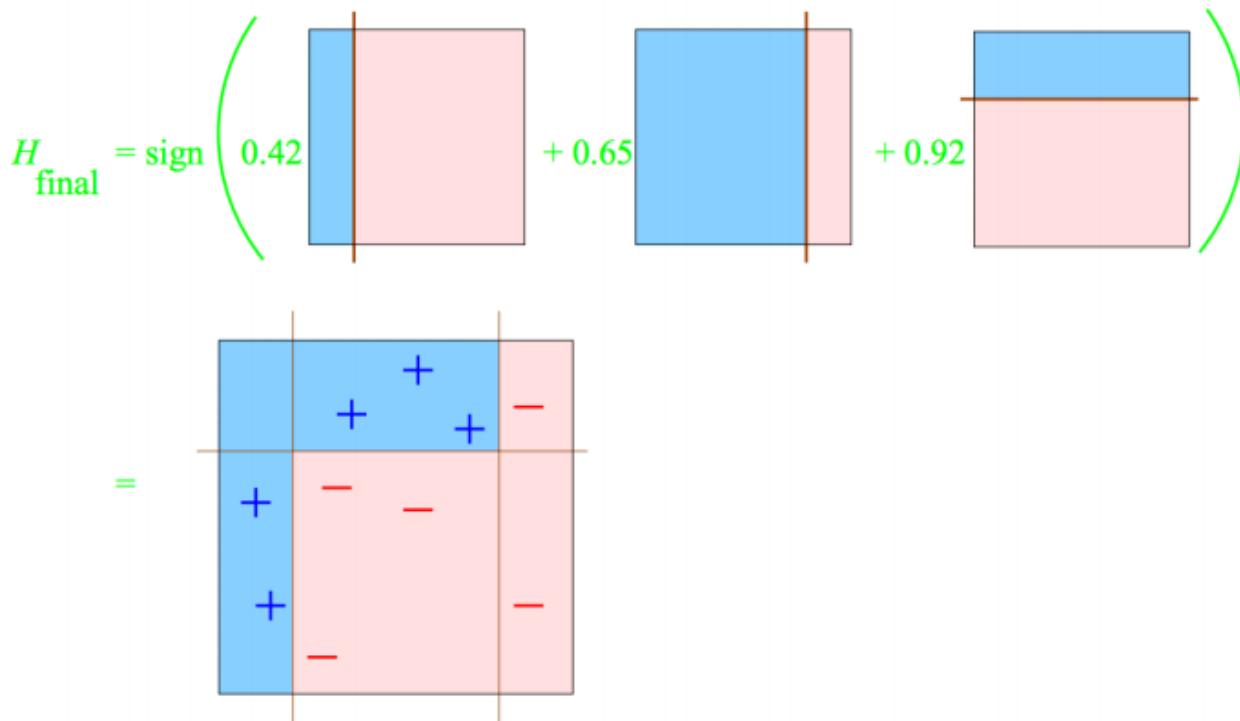


- The 2nd classifier predicts as +/blue everything to the left of the red line (3 samples misclassified)
- The 3rd classifier predicts as +/blue everything above the red line (3 samples misclassified); this classifier tries to correctly predict those samples that were misclassified in round 1 and 2.

Completed classifier: adaboost

- We take the classifiers built in the previous rounds, weight them and add them up

Final Hypothesis



- The resulting classifier is much more accurate! It is the result of the combination of multiple naive simple classifiers (straight lines in this example).

Boosting in R

- Boosting can be used with any subset of (weak) classifiers
- One large subclass is **gradient boosting**
- R has multiple boosting libraries. Differences include the choice of basic classification functions and combination rules:
 - `gbm`: boosting with trees
 - `mboost`: model based boosting
 - `ada`: statistical boosting based on additive logistic regression
 - `gamboost`: for boosting generalised additive models
- Most of these are available in the caret package

Wage example

```
library(ISLR); data(Wage); library(ggplot2); library(caret)
Wage <- subset(Wage, select = -c(logwage)) ## wage is the var want to predict
summary(Wage)
```

```
##      year      age      maritl      race
## Min.   :2003   Min.   :18.00   1. Never Married: 648   1. White:2480
## 1st Qu.:2004   1st Qu.:33.75   2. Married      :2074   2. Black: 293
## Median :2006   Median :42.00   3. Widowed      : 19    3. Asian: 190
## Mean   :2006   Mean   :42.41   4. Divorced     : 204   4. Other:  37
## 3rd Qu.:2008   3rd Qu.:51.00   5. Separated    :  55
## Max.   :2009   Max.   :80.00
##
##      education      region      jobclass
## 1. < HS Grad      :268   2. Middle Atlantic :3000   1. Industrial :1544
## 2. HS Grad        :971   1. New England   :  0    2. Information:1456
## 3. Some College   :650   3. East North Central:  0
## 4. College Grad   :685   4. West North Central:  0
## 5. Advanced Degree:426   5. South Atlantic   :  0
##                      6. East South Central:  0
##                      (Other)      :  0
##      health      health_ins      wage
## 1. <=Good      : 858   1. Yes:2083   Min.   : 20.09
## 2. >=Very Good:2142   2. No : 917   1st Qu.: 85.38
##                      Median :104.92
##                      Mean   :111.70
##                      3rd Qu.:128.68
##                      Max.   :318.34
##
```

```
inTrain<- createDataPartition(y=Wage$wage,p=0.7, list=F)
training<-Wage[inTrain,]; testing<-Wage[-inTrain,]
```


Fit the boosting model

```
modFit<-train(data= training, wage~., method = 'gbm', verbose = F) ##'gbm' to boost with trees// verbose =
F, otherwise too much output is produced
modFit
```

```
## Stochastic Gradient Boosting
##
## 2102 samples
##    9 predictor
##
## No pre-processing
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 2102, 2102, 2102, 2102, 2102, 2102, ...
## Resampling results across tuning parameters:
##
##  interaction.depth  n.trees  RMSE      Rsquared  MAE
##    1                50      33.76863  0.3218465  23.16669
##    1               100      33.25502  0.3325699  22.79275
##    1               150      33.15231  0.3363811  22.76930
##    2                50      33.12180  0.3385880  22.65678
##    2               100      32.99161  0.3428714  22.63563
##    2               150      33.05016  0.3416926  22.71659
##    3                50      33.02860  0.3412073  22.62071
##    3               100      33.10695  0.3400929  22.79829
##    3               150      33.26547  0.3358659  23.00757
##
## Tuning parameter 'shrinkage' was held constant at a value of 0.1
##
## Tuning parameter 'n.minobsinnode' was held constant at a value of 10
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were n.trees = 100, interaction.depth =
## 2, shrinkage = 0.1 and n.minobsinnode = 10.
```

Plot the results

```
qplot(predict(modFit,testing),wage,data=testing)
```



Notes and further reading

- A nice tutorial for boosting: <https://www.cc.gatech.edu/~thad/6601-gradAI-fall2013/boosting.pdf> (<https://www.cc.gatech.edu/~thad/6601-gradAI-fall2013/boosting.pdf>)
- Boosting, random forests and model ensembling are the most common tools that win Kaggle and other prediction contests

Week 3.5: Model-based prediction

Basic ideas

1. Assume the data follow a **probabilistic model**
2. Use the **Bayes' theorem** to identify optimal classifiers based on the probabilistic model previously identified.

Pros: - Can take advantage of structure of the data (e.g. if the data follow a specific probabilistic distribution) - May be computationally convenient - Reasonably accurate on real problems

Cons: - Make **additional assumptions** about the data - When the model is incorrect you may get reduced accuracy

Model based approach

1. Our goal is to build a **parametric model** (a model based on probability distributions) for **conditional distribution** $P(Y = k|X = x)$ (k is a specific y-outcome class)
2. A typical approach is to apply **Bayes theorem** (http://en.wikipedia.org/wiki/Bayes'_theorem):

$$Pr(Y = k|X = x) = \frac{Pr(X = x|Y = k)Pr(Y = k)}{\sum_{\ell=1}^K Pr(X = x|Y = \ell)Pr(Y = \ell)}$$

$$Pr(Y = k|X = x) = \frac{f_k(x)\pi_k}{\sum_{\ell=1}^K f_{\ell}(x)\pi_{\ell}}$$

3. Typically **prior probabilities** π_k are set in advance (using the data).

4. A common choice for $f_k(x) = \frac{1}{\sigma_k \sqrt{2\pi}} e^{-\frac{(x-\mu_k)^2}{\sigma_k^2}}$, a **Gaussian distribution** (might be multivariate gaussian distribution if there are multiple x variables). This is the **parametric model** for the distribution of the features (x) given the outcome class (y=k) we assume.

5. **Estimate the parameters** (μ_k, σ_k^2) from the data.

6. Classify to the **class with the highest value** of $P(Y = k|X = x)$

Classifying using the model

A range of models use this approach

- **Linear discriminant analysis** assumes $f_k(x)$ is **multivariate Gaussian** (i.e. the x features have a multivariate Gaussian distribution within each outcome class) with the **same covariance matrix** for each outcome class. This model draws lines through the data (called the **covariate space**)
- **Quadratic discriminant analysis** assumes $f_k(x)$ is **multivariate Gaussian** with **different covariances** (different covariance matrices are allowed for each outcome class). It draws quadratic curves through the data as opposed to lines.
- **Model based prediction** (<http://www.stat.washington.edu/mclust/>) assumes **more complicated versions for the covariance matrix**
- **Naive Bayes** assumes **independence between features** for model building. This might not be true in reality, although Naive Bayes could still allow to build a useful predictor.

<http://statweb.stanford.edu/~tibs/ElemStatLearn/> (<http://statweb.stanford.edu/~tibs/ElemStatLearn/>)

Why linear discriminant analysis?

$$\log \frac{Pr(Y = k|X = x)}{Pr(Y = j|X = x)}$$

- The **log** is a monotonous function, so when the ratio $\frac{Pr(Y=k|X=x)}{Pr(Y=j|X=x)}$ increases, also $\log \frac{Pr(Y=k|X=x)}{Pr(Y=j|X=x)}$ increases.
- Using Bayes theorem, you can rewrite the first expression as:

$$= \log \frac{f_k(x)}{f_j(x)} + \log \frac{\pi_k}{\pi_j}$$



- The term $\log \frac{f_k(x)}{f_j(x)}$ can be further transformed into two components, the first one depends on the parameters of the Gaussian/normal distributions (i.e. μ_k and μ_j) ...

$$= \log \frac{\pi_k}{\pi_j} - \frac{1}{2} (\mu_k + \mu_j)^T \Sigma^{-1} (\mu_k + \mu_j)$$

...the second one is a **linear term** (for this reason the model draws lines through the data. A variable will have a higher probability to belong to a specific class if it is on one side of the line, and to belong to another class if it is located on the other side of the line):

$$+ x^T \Sigma^{-1} (\mu_k - \mu_j)$$

Decision boundaries



- We are trying to classify samples into 3 categories using 2 predictors (dimensions of the Cartesian axes)
- The Gaussian distributions are visible as ovals in the left panel
- Lines are drawn in the data space when the probability switches from being higher for a class to another class
- The model basically fits Gaussian distributions to the data and uses them to draw lines, classifying the points according to their highest posterior probability

Discriminant function

$$\delta_k(x) = x^T \Sigma^{-1} \mu_k - \frac{1}{2} \mu_k^T \Sigma^{-1} \mu_k + \log(\mu_k)$$

$\Sigma^{-1} \mu_k$ represent the inverse of the covariance matrix for class k (although it is the same for each class in linear discriminant analysis).

$x^T \Sigma^{-1} \mu_k$ is the linear term

- Decide on class based on $\hat{Y}(x) = \text{argmax}_k \delta_k(x)$ (i.e. the value of k which makes the discriminant function bigger for a specific sample)
- We usually estimate parameters with **maximum likelihood**

Naive Bayes

Naive Bayes tries to further simplify the problem. Suppose we have many predictors, we would want to model:

$$P(Y = k | X_1, \dots, X_m)$$

We could use **Bayes Theorem** to get:

$$P(Y = k | X_1, \dots, X_m) = \frac{\pi_k P(X_1, \dots, X_m | Y = k)}{\sum_{\ell=1}^K P(X_1, \dots, X_m | Y = \ell) \pi_\ell}$$

It can be claimed that $P(Y = k | X_1, \dots, X_m)$ is **proportional** to the numerator of Bayes theorem (the term in the denominator is just a constant for all the different probabilities):

$$\propto \pi_k P(X_1, \dots, X_m | Y = k)$$

This can be re-written (breaking down **conditional probabilities** until you get one term for every feature X):

$$P(X_1, \dots, X_m, Y = k) = \pi_k P(X_1 | Y = k) P(X_2, \dots, X_m | X_1, Y = k)$$

$$= \pi_k P(X_1|Y = k)P(X_2|X_1, Y = k)P(X_3, \dots, X_m|X_1, X_2, Y = k)$$

$$= \pi_k P(X_1|Y = k)P(X_2|X_1, Y = k) \dots P(X_m|X_1 \dots, X_{m-1}, Y = k)$$

We could make an **assumption (if all X s were independent)** to write this (where the conditioning is dropped down):

$$\approx \pi_k P(X_1|Y = k)P(X_2|Y = k) \dots P(X_m|Y = k)$$

- Although this assumption is **naive**, it works particularly well in a number of applications (e.g. when you have a very large of binary/categorical features, for example in text/document classification applications)

Example: Iris Data

```
data(iris); library(ggplot2)
names(iris)
```

```
## [1] "Sepal.Length" "Sepal.Width" "Petal.Length" "Petal.Width" "Species"
```

```
table(iris$Species)
```

```
##
##      setosa versicolor virginica
##         50         50         50
```

Create training and test sets

```
library(caret)
inTrain<-createDataPartition(y=iris$Species, p=0.7,list = F)
training<-iris[inTrain,]
testing<-iris[-inTrain,]
dim(training); dim(testing)
```

```
## [1] 105  5
```

```
## [1] 45  5
```


Build predictions

```
modLDA<-train(data=training, Species~., method='lda')
modNB<-train(data=training, Species~., method='nb')
pdLDA<- predict(modLDA, testing)
pdNB<- predict(modNB, testing)
table(pdLDA,pdNB)
```

```
##
##      pdLDA      pdNB
##      setosa versicolor virginica
##      setosa      15         0         0
##      versicolor   0        13         2
##      virginica    0         1        14
```

Comparison of results

```
equalPredictions <- (pdLDA == pdNB)
qplot(Petal.Width, Sepal.Width, colour=equalPredictions, data=testing)
```



Notes and further reading

- Introduction to statistical learning (<http://www-bcf.usc.edu/~gareth/ISL/>)
- Elements of Statistical Learning (<http://www-stat.stanford.edu/~tibs/ElemStatLearn/>)
- Model based clustering (<http://www.stat.washington.edu/raftery/Research/PDF/fraley2002.pdf>)
- Linear Discriminant Analysis (http://en.wikipedia.org/wiki/Linear_discriminant_analysis)
- Quadratic Discriminant Analysis (http://en.wikipedia.org/wiki/Quadratic_classifier)